

# CS156 (Introduction to AI), Spring 2021

## Homework 9 submission

Roster Name: Neeval Kumar

Preferred Name (if different): Chosen Name

Student ID: 011877086

Email address: kumar.neeval@gmail.com

Any special notes or anything you would like to communicate to me about this homework submission goes in here.

## References and sources

List all your references and sources here. This includes all sites/discussion boards/blogs/posts/etc. where you grabbed some code examples.

I only used tensorflow's website for all the libraries.

<https://www.tensorflow.org/tutorials>

## Solution

Load libraries and set random number generator seed

```
In [1]: import tensorflow as tf
        from tensorflow import keras
        from tensorflow.keras import layers
        import os
        import matplotlib.pyplot as plt
        from skimage import io
        import numpy as np
        from sklearn.metrics import plot_confusion_matrix
        import seaborn as sns
        from sklearn.metrics import accuracy_score
```

```
In [2]: np.random.seed(42)
```

## Load images into keras dataset

```
In [3]: image_size = (180, 180)
batch_size = 32

print("Preprocessing Training set: ")
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "flowers/training",
    labels='inferred',
    label_mode='categorical',
    validation_split=0.2,
    subset="training",
    seed=42,
    image_size=image_size,
    batch_size=batch_size,)

print("\nPreprocessing Validation set: ")
val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "flowers/training",
    validation_split=0.2,
    labels='inferred',
    label_mode='categorical',
    subset="validation",
    seed=42,
    image_size=image_size,
    batch_size=batch_size,
)

class_names = train_ds.class_names

print("\nClass labels are: \n", class_names)
```

```
Preprocessing Training set:
Found 3456 files belonging to 5 classes.
Using 2765 files for training.
```

```
Preprocessing Validation set:
Found 3456 files belonging to 5 classes.
Using 691 files for validation.
```

```
Class labels are:
['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']
```


## Plot some images with their true labels

```
In [4]: plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        arr = labels[i].numpy()
        index_of = np.where(arr==1.)
        index_of = int(index_of[0])
        plt.title(class_names[index_of])
        plt.axis("off")
```

tulip



sunflower



daisy



sunflower




tulip




rose




tulip



dandelion






tulip



Augment Images and then plot augments

```
In [5]: data_augmentation = keras.Sequential(  
    [  
        layers.experimental.preprocessing.RandomFlip("horizontal"),  
        layers.experimental.preprocessing.RandomRotation(0.1),  
    ]  
)
```

```
In [6]: plt.figure(figsize=(10, 10))  
for images, _ in train_ds.take(1):  
    for i in range(9):  
        augmented_images = data_augmentation(images)  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(augmented_images[0].numpy().astype("uint8"))  
        plt.axis("off")
```



CNN that makes a model



```

In [7]: def make_model(input_shape, num_classes):
    inputs = keras.Input(shape=input_shape)
    # Image augmentation block
    x = data_augmentation(inputs)

    # Entry block
    x = layers.experimental.preprocessing.Rescaling(1.0 / 255)(x)
    x = layers.Conv2D(32, 3, strides=2, padding="same")(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation("relu")(x)

    x = layers.Conv2D(64, 3, padding="same")(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation("relu")(x)

    previous_block_activation = x # Set aside residual

    for size in [128, 256, 512, 728]:
        x = layers.Activation("relu")(x)
        x = layers.SeparableConv2D(size, 3, padding="same")(x)
        x = layers.BatchNormalization()(x)

        x = layers.Activation("relu")(x)
        x = layers.SeparableConv2D(size, 3, padding="same")(x)
        x = layers.BatchNormalization()(x)

        x = layers.MaxPooling2D(3, strides=2, padding="same")(x)

        # Project residual
        residual = layers.Conv2D(size, 1, strides=2, padding="same")(
            previous_block_activation
        )
        x = layers.add([x, residual]) # Add back residual
        previous_block_activation = x # Set aside next residual

    x = layers.SeparableConv2D(1024, 3, padding="same")(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation("relu")(x)

    x = layers.GlobalAveragePooling2D()(x)
    if num_classes == 2:
        activation = "sigmoid"
        units = 1
    else:
        activation = "softmax"
        units = num_classes

    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(units, activation=activation)(x)
    return keras.Model(inputs, outputs)

```

```
In [8]: train_ds = train_ds.prefetch(buffer_size=32)
val_ds = val_ds.prefetch(buffer_size=32)
model = make_model(input_shape=image_size + (3,), num_classes=5)
#keras.utils.plot_model(model, show_shapes=True)
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 180, 180, 3)]	0	
sequential (Sequential)	(None, 180, 180, 3)	0	input_1[0][0]
rescaling (Rescaling) [0]	(None, 180, 180, 3)	0	sequential[0]
conv2d (Conv2D) [0]	(None, 90, 90, 32)	896	rescaling[0][0]
batch_normalization (BatchNorma	(None, 90, 90, 32)	128	conv2d[0][0]
activation (Activation) zation[0][0]	(None, 90, 90, 32)	0	batch_normali
conv2d_1 (Conv2D) [0]	(None, 90, 90, 64)	18496	activation[0]
batch_normalization_1 (BatchNor	(None, 90, 90, 64)	256	conv2d_1[0][0]
activation_1 (Activation) zation_1[0][0]	(None, 90, 90, 64)	0	batch_normali
activation_2 (Activation) [0][0]	(None, 90, 90, 64)	0	activation_1[0][0]
separable_conv2d (SeparableConv	(None, 90, 90, 128)	8896	activation_2[0][0]
batch_normalization_2 (BatchNor	(None, 90, 90, 128)	512	separable_con

activation_3 (Activation) zation_2[0][0]	(None, 90, 90, 128)	0	batch_normali
separable_conv2d_1 (SeparableCo 0][0]	(None, 90, 90, 128)	17664	activation_3[
batch_normalization_3 (BatchNor v2d_1[0][0]	(None, 90, 90, 128)	512	separable_con
max_pooling2d (MaxPooling2D) zation_3[0][0]	(None, 45, 45, 128)	0	batch_normali
conv2d_2 (Conv2D) 0][0]	(None, 45, 45, 128)	8320	activation_1[
add (Add) 0][0]	(None, 45, 45, 128)	0	max_pooling2d conv2d_2[0][0]
activation_4 (Activation)	(None, 45, 45, 128)	0	add[0][0]
separable_conv2d_2 (SeparableCo 0][0]	(None, 45, 45, 256)	34176	activation_4[
batch_normalization_4 (BatchNor v2d_2[0][0]	(None, 45, 45, 256)	1024	separable_con
activation_5 (Activation) zation_4[0][0]	(None, 45, 45, 256)	0	batch_normali
separable_conv2d_3 (SeparableCo 0][0]	(None, 45, 45, 256)	68096	activation_5[
batch_normalization_5 (BatchNor v2d_3[0][0]	(None, 45, 45, 256)	1024	separable_con
max_pooling2d_1 (MaxPooling2D) zation_5[0][0]	(None, 23, 23, 256)	0	batch_normali
conv2d_3 (Conv2D)	(None, 23, 23, 256)	33024	add[0][0]

add_1 (Add) _1[0][0]	(None, 23, 23, 256)	0	max_pooling2d conv2d_3[0][0]
<hr/>			
activation_6 (Activation)	(None, 23, 23, 256)	0	add_1[0][0]
<hr/>			
separable_conv2d_4 (SeparableCo v2d_4[0][0])	(None, 23, 23, 512)	133888	activation_6[0][0]
<hr/>			
batch_normalization_6 (BatchNor v2d_4[0][0])	(None, 23, 23, 512)	2048	separable_con
<hr/>			
activation_7 (Activation) zation_6[0][0]	(None, 23, 23, 512)	0	batch_normali
<hr/>			
separable_conv2d_5 (SeparableCo v2d_5[0][0])	(None, 23, 23, 512)	267264	activation_7[0][0]
<hr/>			
batch_normalization_7 (BatchNor v2d_5[0][0])	(None, 23, 23, 512)	2048	separable_con
<hr/>			
max_pooling2d_2 (MaxPooling2D) zation_7[0][0]	(None, 12, 12, 512)	0	batch_normali
<hr/>			
conv2d_4 (Conv2D)	(None, 12, 12, 512)	131584	add_1[0][0]
<hr/>			
add_2 (Add) _2[0][0]	(None, 12, 12, 512)	0	max_pooling2d conv2d_4[0][0]
<hr/>			
activation_8 (Activation)	(None, 12, 12, 512)	0	add_2[0][0]
<hr/>			
separable_conv2d_6 (SeparableCo v2d_6[0][0])	(None, 12, 12, 728)	378072	activation_8[0][0]
<hr/>			
batch_normalization_8 (BatchNor v2d_6[0][0])	(None, 12, 12, 728)	2912	separable_con
<hr/>			
activation_9 (Activation) zation_8[0][0]	(None, 12, 12, 728)	0	batch_normali
<hr/>			



separable_conv2d_7 (SeparableCo	(None, 12, 12, 728)	537264	activation_9[0][0]
batch_normalization_9 (BatchNor	(None, 12, 12, 728)	2912	separable_conv2d_7[0][0]
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 728)	0	batch_normalization_9[0][0]
conv2d_5 (Conv2D)	(None, 6, 6, 728)	373464	add_2[0][0]
add_3 (Add)	(None, 6, 6, 728)	0	max_pooling2d_3[0][0]
separable_conv2d_8 (SeparableCo	(None, 6, 6, 1024)	753048	add_3[0][0]
batch_normalization_10 (BatchNo	(None, 6, 6, 1024)	4096	separable_conv2d_8[0][0]
activation_10 (Activation)	(None, 6, 6, 1024)	0	batch_normalization_10[0][0]
global_average_pooling2d (Globa	(None, 1024)	0	activation_10[0][0]
dropout (Dropout)	(None, 1024)	0	global_average_pooling2d[0][0]
dense (Dense)	(None, 5)	5125	dropout[0][0]
=====			
Total params: 2,786,749			
Trainable params: 2,778,013			
Non-trainable params: 8,736			

## Train Model

```
In [9]: epochs = 20

callbacks = [
    keras.callbacks.ModelCheckpoint("save_at_{epoch}.h5"),
]
model.compile(
    optimizer=keras.optimizers.Adam(1e-3),
    loss="categorical_crossentropy",
    metrics=["accuracy"],
)
model.fit(
    train_ds, epochs=epochs, callbacks=callbacks, validation_data=val_ds,
)
```

```
Epoch 1/20
87/87 [=====] - 356s 4s/step - loss: 1.3189 - accuracy: 0.4968 - val_loss: 1.7162 - val_accuracy: 0.2590
Epoch 2/20
87/87 [=====] - 387s 4s/step - loss: 1.0320 - accuracy: 0.6204 - val_loss: 2.2801 - val_accuracy: 0.2590
Epoch 3/20
87/87 [=====] - 341s 4s/step - loss: 0.9433 - accuracy: 0.6565 - val_loss: 3.1217 - val_accuracy: 0.2590
Epoch 4/20
87/87 [=====] - 339s 4s/step - loss: 0.8775 - accuracy: 0.6830 - val_loss: 3.2253 - val_accuracy: 0.2590
Epoch 5/20
87/87 [=====] - 6634s 77s/step - loss: 0.7986 - accuracy: 0.6966 - val_loss: 3.7676 - val_accuracy: 0.2590
Epoch 6/20
87/87 [=====] - 4233s 49s/step - loss: 0.7009 - accuracy: 0.7308 - val_loss: 2.0200 - val_accuracy: 0.3415
Epoch 7/20
87/87 [=====] - 361s 4s/step - loss: 0.7344 - accuracy: 0.7372 - val_loss: 1.0680 - val_accuracy: 0.6049
Epoch 8/20
87/87 [=====] - 410s 5s/step - loss: 0.6332 - accuracy: 0.7642 - val_loss: 0.9283 - val_accuracy: 0.6787
Epoch 9/20
87/87 [=====] - 389s 4s/step - loss: 0.6122 - accuracy: 0.7811 - val_loss: 0.7668 - val_accuracy: 0.7352
Epoch 10/20
87/87 [=====] - 404s 5s/step - loss: 0.6197 - accuracy: 0.7649 - val_loss: 1.9760 - val_accuracy: 0.4964
Epoch 11/20
87/87 [=====] - 407s 5s/step - loss: 0.5665 - accuracy: 0.7918 - val_loss: 0.7083 - val_accuracy: 0.7410
Epoch 12/20
87/87 [=====] - 401s 5s/step - loss: 0.5752 - accuracy: 0.7831 - val_loss: 0.6905 - val_accuracy: 0.7757
Epoch 13/20
87/87 [=====] - 408s 5s/step - loss: 0.5498 - accuracy: 0.7880 - val_loss: 0.9262 - val_accuracy: 0.7004
Epoch 14/20
87/87 [=====] - 416s 5s/step - loss: 0.5377 - accuracy:
```

```

y: 0.7976 - val_loss: 0.7034 - val_accuracy: 0.7931
Epoch 15/20
87/87 [=====] - 407s 5s/step - loss: 0.4841 - accurac
y: 0.8209 - val_loss: 0.5964 - val_accuracy: 0.8061
Epoch 16/20
87/87 [=====] - 712s 8s/step - loss: 0.4397 - accurac
y: 0.8320 - val_loss: 1.1332 - val_accuracy: 0.7091
Epoch 17/20
87/87 [=====] - 358s 4s/step - loss: 0.4976 - accurac
y: 0.8151 - val_loss: 0.6495 - val_accuracy: 0.7815
Epoch 18/20
87/87 [=====] - 358s 4s/step - loss: 0.4395 - accurac
y: 0.8324 - val_loss: 0.6946 - val_accuracy: 0.7887
Epoch 19/20
87/87 [=====] - 383s 4s/step - loss: 0.4417 - accurac
y: 0.8330 - val_loss: 0.5675 - val_accuracy: 0.8148
Epoch 20/20
87/87 [=====] - 374s 4s/step - loss: 0.4256 - accurac
y: 0.8451 - val_loss: 0.6091 - val_accuracy: 0.8191

```

Out[9]: <tensorflow.python.keras.callbacks.History at 0x7fa814a18af0>

## Load Test Dataset

```

In [10]: test_ds = tf.keras.preprocessing.image_dataset_from_directory(
            "flowers/test",
            labels='inferred',
            label_mode='categorical',
            seed=42,
            image_size=image_size,
            batch_size=1,
        )

```

Found 861 files belonging to 5 classes.

**Predict images from test set onto model, create lists for true labels and predicted labels. Also store 3 misclassified for visualization**

```
In [11]: true_labels = []
predicted_labels = []
misclassified = []
count = 0
for images, label in test_ds:
    predictions = model.predict(images)
    best_prediction = np.argmax(predictions)
    true_label = np.argmax(label)
    predicted_labels.append(best_prediction)
    true_labels.append(true_label)

    if best_prediction != true_label and count != 3:
        misclass_arr = [label, images, best_prediction, true_label]
        misclassified.append(misclass_arr)
        count += 1

print("Finished test set onto model")
```

Finished test set onto model

**Print Accuracy score of true labels and predicted labels, then plot confusion matrix**


```
In [12]: print("Accuracy score: ", accuracy_score(true_labels, predicted_labels))

cm = tf.math.confusion_matrix(labels=true_labels, predictions=predicted_labels)

sns.heatmap(
    cm, annot=True, fmt = 'd',
    xticklabels=class_names,
    yticklabels=class_names)
plt.xlabel("Predicted")
plt.ylabel("True")
```

Accuracy score: 0.7073170731707317

Out[12]: Text(33.0, 0.5, 'True')



## Display 3 misclassified Images

```
In [13]: plt.figure(figsize=(5, 5))
for x in misclassified:
    fig, ax = plt.subplots(nrows=1, ncols=1)
    label = x[0]
    new_image = x[1]
    img_title = str(class_names[x[3]]) + " predicted as " + str(class_names[
    plt.imshow(new_image[0].numpy().astype("uint8"))
    plt.title(img_title)
    plt.axis("off")
```

<Figure size 360x360 with 0 Axes>

sunflower predicted as dandelion



sunflower predicted as daisy



dandelion predicted as daisy



In [ ]: